

# Project\_BigData

December 11, 2019

```
[89]: from pyspark.sql import functions as F
      from pyspark.sql import DataFrameNaFunctions as DFna
      from pyspark.sql.functions import udf, col, when
      import matplotlib.pyplot as plt
      import pyspark as ps
      import os, sys, requests, json
      from pyspark.ml.evaluation import RegressionEvaluator
      from pyspark.ml.recommendation import ALS
      from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
      from pyspark.ml import Pipeline
      from pyspark.sql import Row
      import numpy as np
      import math
      import pandas as pd
      from pandas import Series, DataFrame
```

```
[90]: spark = ps.sql.SparkSession.builder \
      .master("local[4]") \
      .appName("building recommender") \
      .getOrCreate() # create a spark session

sc = spark.sparkContext
```

```
[36]: # read movies CSV
      movies = spark.read.option("header", "true").csv("/Users/vkoushikmuthyapu/
      ↳desktop/ml-latest-small/movies.csv",inferSchema=True)
      movies.printSchema()
      movies.show()
```

```
root
 |-- movieId: integer (nullable = true)
 |-- title: string (nullable = true)
 |-- genres: string (nullable = true)

+-----+-----+-----+
|movieId|          title|          genres|
+-----+-----+-----+
|      1| Toy Story (1995)|Adventure|Animati...
```

2	Jumanji (1995)	Adventure	Childre...
3	Grumpier Old Men ...	Comedy	Romance
4	Waiting to Exhale...	Comedy	Drama Romance
5	Father of the Bri...	Comedy	
6	Heat (1995)	Action	Crime Thri...
7	Sabrina (1995)	Comedy	Romance
8	Tom and Huck (1995)	Adventure	Children
9	Sudden Death (1995)	Action	
10	GoldenEye (1995)	Action	Adventure ...
11	American Presiden...	Comedy	Drama Romance
12	Dracula: Dead and...	Comedy	Horror
13	Balto (1995)	Adventure	Animati...
14	Nixon (1995)	Drama	
15	Cutthroat Island ...	Action	Adventure ...
16	Casino (1995)	Crime	Drama
17	Sense and Sensibi...	Drama	Romance
18	Four Rooms (1995)	Comedy	
19	Ace Ventura: When...	Comedy	
20	Money Train (1995)	Action	Comedy Cri...

+-----+-----+-----+-----+

only showing top 20 rows

```
[37]: ratings = spark.read.option("header", "true").csv("/Users/vkoushikmuthyapu/
↳desktop/ml-latest-small/ratings.csv",inferSchema=True)
ratings.printSchema()
ratings.show()
```

root

```
|-- userId: integer (nullable = true)
|-- movieId: integer (nullable = true)
|-- rating: double (nullable = true)
|-- timestamp: integer (nullable = true)
```

userId	movieId	rating	timestamp
1	1	4.0	964982703
1	3	4.0	964981247
1	6	4.0	964982224
1	47	5.0	964983815
1	50	5.0	964982931
1	70	3.0	964982400
1	101	5.0	964980868
1	110	4.0	964982176
1	151	5.0	964984041
1	157	5.0	964984100
1	163	5.0	964983650

	1	216	5.0 964981208
	1	223	3.0 964980985
	1	231	5.0 964981179
	1	235	4.0 964980908
	1	260	5.0 964981680
	1	296	3.0 964982967
	1	316	3.0 964982310
	1	333	5.0 964981179
	1	349	4.0 964982563

+-----+-----+-----+-----+

only showing top 20 rows

```
[38]: newrating = ratings.select(['userId', 'movieId', 'rating'])
      newrating.show()
```

+-----+-----+-----+

	userId	movieId	rating
--	--------	---------	--------

+-----+-----+-----+

	1	1	4.0
	1	3	4.0
	1	6	4.0
	1	47	5.0
	1	50	5.0
	1	70	3.0
	1	101	5.0
	1	110	4.0
	1	151	5.0
	1	157	5.0
	1	163	5.0
	1	216	5.0
	1	223	3.0
	1	231	5.0
	1	235	4.0
	1	260	5.0
	1	296	3.0
	1	316	3.0
	1	333	5.0
	1	349	4.0

+-----+-----+-----+

only showing top 20 rows

```
[39]: newerratings = newrating.rdd
      newerratings
```

```
[39]: MapPartitionsRDD[706] at javaToPython at NativeMethodAccessorImpl.java:0
```

```
[40]: training_df, validation_df, test_df = newrating.randomSplit([.6, .2, .2],
    ↪seed=0)
#training_RDD = training_df.rdd.map(lambda x: (x[0], x[1])).cache()
#validation_for_predict_RDD = validation_df.rdd
#test_for_predict_RDD = test_df.rdd.map(lambda x: (x[0], x[1])).cache()
training_df
```

```
[40]: DataFrame[userId: int, movieId: int, rating: double]
```

```
[132]: als = ALS(maxIter=10, regParam=0.05, rank=18, userCol="userId",
    ↪itemCol="movieId", ratingCol="rating", coldStartStrategy="drop",
    ↪nonnegative= True)          # regularization param)
model = als.fit(training_df)
# make prediction
predictions = model.transform(validation_df)
new_predictions = predictions.filter(col('prediction') != np.nan)
rmse = evaluator.evaluate(new_predictions)
print ("For rank =",18, "reg =", 0.05 , " the RMSE= " ,rmse)
```

For rank = 18 reg = 0.05 the RMSE= 0.9803222190248909

```
[41]: iterations = 10
regularization_parameter = [0.001, 0.01, 0.05, 0.1, 0.2]
ranks = [8, 10, 12, 14, 16, 18, 20]
errors = []
err = 0
#tolerance = 0.02
```

```
[80]: min_error = float('inf')
best_rank = -1
best_iteration = -1

for rank in ranks:
    for reg in regularization_parameter:
        # train ALS model
        als = ALS(maxIter=iterations, regParam=reg, rank=rank,
    ↪userCol="userId", itemCol="movieId", ratingCol="rating",
    ↪coldStartStrategy="drop", nonnegative= True)          # regularization param)
        #model = als.fit(training_df)
        # make prediction
        #predictions = model.transform(validation_df)
        #new_predictions = predictions.filter(col('prediction') != np.nan)
        param_grid= ParamGridBuilder().addGrid(als.rank,[rank]).addGrid(als.
    ↪maxIter,[10]).addGrid(als.regParam,[reg]).build()
        evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
    ↪predictionCol="prediction")
        crossval = CrossValidator(estimator=als,
```

```

        estimatorParamMaps=param_grid,
        evaluator=evaluator,
        numFolds=5)
cvModel = crossval.fit(training_df)
cvModel_pred = cvModel.transform(validation_df)
cvModel_pred = cvModel_pred.filter(col('prediction') != np.nan)
rmse = evaluator.evaluate(cvModel_pred)
errors.append(rmse)

print ("For rank =",rank, "reg =", reg , " the RMSE= " ,rmse)
if rmse < min_error:
    min_error = rmse
    best_rank = rank
    best_reg = reg
print ("The best model was trained with rank= ", best_rank, "With reg= ",
↪best_reg)

```

```

For rank = 8 reg = 0.001 the RMSE= 1.4467448634372828
For rank = 8 reg = 0.01 the RMSE= 1.141916955112322
For rank = 8 reg = 0.05 the RMSE= 0.9856996845615988
For rank = 8 reg = 0.1 the RMSE= 0.9141869244440859
For rank = 8 reg = 0.2 the RMSE= 0.8966076281249638
For rank = 10 reg = 0.001 the RMSE= 1.456039625308575
For rank = 10 reg = 0.01 the RMSE= 1.2036612472694024
For rank = 10 reg = 0.05 the RMSE= 0.9956121741381565
For rank = 10 reg = 0.1 the RMSE= 0.9158487323974703
For rank = 10 reg = 0.2 the RMSE= 0.8964889415116521
For rank = 12 reg = 0.001 the RMSE= 1.5242743892266397
For rank = 12 reg = 0.01 the RMSE= 1.235516337012286
For rank = 12 reg = 0.05 the RMSE= 1.0054946510104066
For rank = 12 reg = 0.1 the RMSE= 0.9178090013903915
For rank = 12 reg = 0.2 the RMSE= 0.8985851115100361
For rank = 14 reg = 0.001 the RMSE= 1.5311798036079438
For rank = 14 reg = 0.01 the RMSE= 1.2574262798813995
For rank = 14 reg = 0.05 the RMSE= 1.004432265748177
For rank = 14 reg = 0.1 the RMSE= 0.9163648847902192
For rank = 14 reg = 0.2 the RMSE= 0.8990094563068959
For rank = 16 reg = 0.001 the RMSE= 1.5977755203925776
For rank = 16 reg = 0.01 the RMSE= 1.2945525589067843
For rank = 16 reg = 0.05 the RMSE= 1.004848074487897
For rank = 16 reg = 0.1 the RMSE= 0.9149908997725777
For rank = 16 reg = 0.2 the RMSE= 0.8988097775919107
For rank = 18 reg = 0.001 the RMSE= 1.64173527955219
For rank = 18 reg = 0.01 the RMSE= 1.3049537155223991
For rank = 18 reg = 0.05 the RMSE= 1.008261040265912
For rank = 18 reg = 0.1 the RMSE= 0.915377249956471

```

```

For rank = 18 reg = 0.2 the RMSE= 0.8988576731676955
For rank = 20 reg = 0.001 the RMSE= 1.7042923866197208
For rank = 20 reg = 0.01 the RMSE= 1.3538113042962867
For rank = 20 reg = 0.05 the RMSE= 1.0154965096385316
For rank = 20 reg = 0.1 the RMSE= 0.9160800093109257
For rank = 20 reg = 0.2 the RMSE= 0.8999188491754373
The best model was trained with rank= 10 With reg= 0.2

```

```

└─
↪-----

NameError                                Traceback (most recent call
↪last)

<ipython-input-80-738ecb0fa263> in <module>
    30             best_reg = reg
    31 print ("The best model was trained with rank= ", best_rank, "With
↪reg= ", best_reg)
---> 32 best = model.bestModel

NameError: name 'model' is not defined

```

```
[82]: best = cvModel.bestModel
```

```
best = ALS_42b8a21f89f7ed1ddbdf
```

```
[102]: #here testing with new test data
als = ALS(maxIter=iterations, regParam=0.2, rank=10, userCol="userId",
↪itemCol="movieId", ratingCol="rating")
#param_grid= ParamGridBuilder().addGrid(als.rank,[10]).addGrid(als.
↪maxIter,[10]).addGrid(als.regParam,[0.2]).build()
cvModel = als.fit(training_df)
evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
↪predictionCol="prediction")
cvModel_pred = cvModel.transform(test_df)
cvModel_pred = cvModel_pred.filter(col('prediction') != np.nan)
rmseT = evaluator.evaluate(cvModel_pred)
print("test data Rmse= ", rmseT)
#display(cvModel_pred.sort("userID", "ratings"))

```

```
test data Rmse= 0.8935807766860681
```

```
[103]: prediction = (cvModel_pred.sort(newrating["userID"]))
prediction.show()
```

userId	movieId	rating	prediction
1	2143	4.0	3.4857223
1	2959	5.0	4.8959446
1	736	3.0	3.546459
1	3273	5.0	2.521188
1	2078	5.0	4.5469894
1	553	5.0	4.2281413
1	733	4.0	4.186975
1	2141	5.0	3.6555433
1	423	3.0	3.2964041
1	2654	5.0	3.6944335
1	1270	5.0	4.675041
1	47	5.0	4.542998
1	1136	5.0	4.7466288
1	1226	5.0	3.6889038
1	3527	4.0	4.1616406
1	367	4.0	3.7773545
1	1	4.0	4.5714073
1	500	3.0	3.9910703
1	2987	5.0	4.081792
1	1025	5.0	4.2760987

only showing top 20 rows

```
[104]: cvModel_pred = cvModel_pred.na.drop()
cvModel_pred.describe().show()
```

summary	userId	movieId	rating
count	19140	19140	19140
mean	322.9807732497388	17542.324503657263	3.5158307210031348
stddev	181.02445275760041	32812.525669396804	1.0359594289983285
min	1	1	0.5
max	610	189333	5.0

```
[125]: user_recs = best.recommendForAllUsers(10)
       user_recs
```

```
[125]: DataFrame[userId: int, recommendations: array<struct<movieId:int,rating:float>>]
```


```
[126]: def recs_users(recs):
       recs = recs.select("recommendations.movieId", "recommendations.rating")
       movies = recs.select("movieId").toPandas().iloc[0,0]
       ratings = recs.select("rating").toPandas().iloc[0,0]
       ratings_matrix = pd.DataFrame(movies, columns = ["movieId"])
       ratings_matrix["ratings"] = ratings
       ratings_matrix_ps = ratings_matrix
       return ratings_matrix_ps
```

```
[127]: test = recs_users(user_recs)
       test.join(movies)
```


```
[127]:   movieId  ratings   movieId,title,genres
0      3379  4.701992  1,Toy Story (1995),Adventure|Animation|Childre...
1      6818  4.486014      2,Jumanji (1995),Adventure|Children|Fantasy
2      3358  4.478576      3,Grumpier Old Men (1995),Comedy|Romance
3      5915  4.441676      4,Waiting to Exhale (1995),Comedy|Drama|Romance
4      5490  4.441676      5,Father of the Bride Part II (1995),Comedy
5      99764  4.402958      6,Heat (1995),Action|Crime|Thriller
6     148881  4.402958      7,Sabrina (1995),Comedy|Romance
7      40491  4.402958      8,Tom and Huck (1995),Adventure|Children
8       8477  4.402958      9,Sudden Death (1995),Action
9       3153  4.383376     10,GoldenEye (1995),Action|Adventure|Thriller
```

```
[135]: cvModel_pred.toPandas()['rating'].hist()
       plt.xlabel('ratings')
       plt.ylabel('users')
       plt.title('rating vs users')
       plt.show()
```





```
[136]: cvModel_pred.toPandas()['prediction'].hist()  
plt.xlabel('predicted_ratings')  
plt.ylabel('users')  
plt.title('predicted_ratings vs users')  
plt.show()
```



[ ]:

[ ]: